

SYSU 18364012 陈鑫锐

OS ASSIGNMENT3

1. Xv6 lab: Multithreading/Uthread: switching between threads

第一步：在 `uthread_switch.S` 中实现 `thread_switch` 函数，参考 `kernel/switch.S`，保存当前线程的寄存器，恢复即将要切换到线程的寄存器

```
thread_switch:
    /* YOUR CODE HERE */

    sd ra, 0(a0)
    sd sp, 8(a0)
    sd s0, 16(a0)
    sd s1, 24(a0)
    sd s2, 32(a0)
    sd s3, 40(a0)
    sd s4, 48(a0)
    sd s5, 56(a0)
    sd s6, 64(a0)
    sd s7, 72(a0)
    sd s8, 80(a0)
    sd s9, 88(a0)
    sd s10, 96(a0)
    sd s11, 104(a0)

    ld ra, 0(a1)
    ld sp, 8(a1)
    ld s0, 16(a1)
    ld s1, 24(a1)
    ld s2, 32(a1)
    ld s3, 40(a1)
    ld s4, 48(a1)
    ld s5, 56(a1)
    ld s6, 64(a1)
    ld s7, 72(a1)
    ld s8, 80(a1)
    ld s9, 88(a1)
    ld s10, 96(a1)
    ld s11, 104(a1)
    ret
```

第二步：在 `uthread.c` 的 `struct thread` 添加寄存器字段

```
struct thread {
    /*stored registers*/
    uint64 ra;
    uint64 sp;

    // callee-saved
    uint64 s0;
    uint64 s1;
    uint64 s2;
    uint64 s3;
    uint64 s4;
    uint64 s5;
    uint64 s6;
    uint64 s7;
    uint64 s8;
    uint64 s9;
    uint64 s10;
    uint64 s11;

    char    stack[STACK_SIZE]; /* the thread's stack */
    int     state;              /* FREE, RUNNING, RUNNABLE */
};
```

第三步：在 `thread_create()` 中添加保存新线程返回地址和栈指针，返回地址就是输入的函数指针 `func`，栈指针指向 `struct thread->stack` 的最后一个元素的地址，因为栈指针是从高地址向低地址增长的

```
void
thread_create(void (*func)())
{
    struct thread *t;

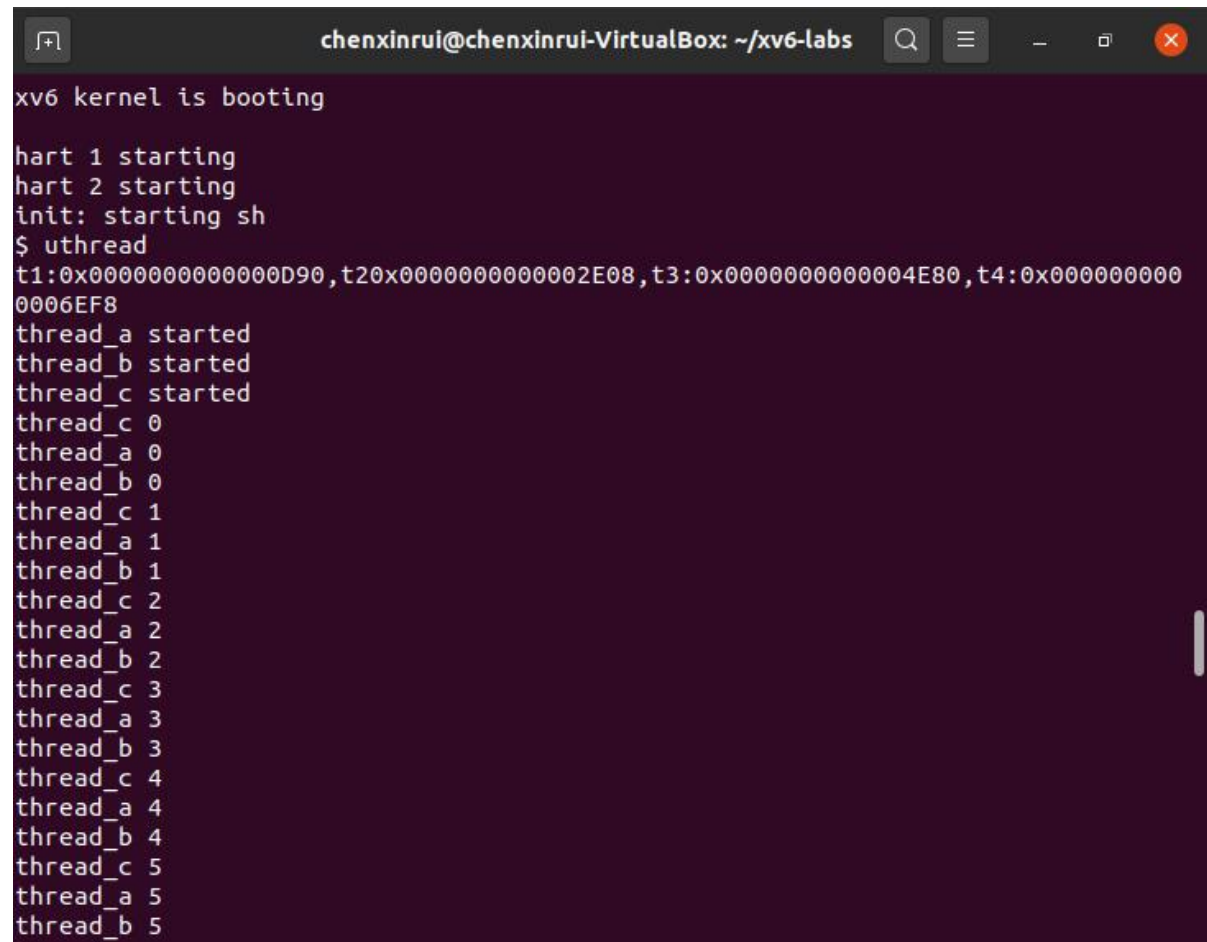
    for (t = all_thread; t < all_thread + MAX_THREAD; t++) {
        if (t->state == FREE) break;
    }
    t->state = RUNNABLE;
    // YOUR CODE HERE

    t->ra=(uint64)func;
    t->sp=(uint64)&t->stack[STACK_SIZE-1];
}
```

第四步：最后在 `thread_schedule()` 中添加调用

```
/* YOUR CODE HERE
 * Invoke thread_switch to switch from t to next_thread:*/
thread_switch((uint64)t,(uint64)next_thread);
```

运行结果：



```
chenxinrui@chenxinrui-VirtualBox: ~/xv6-labs
xv6 kernel is booting
hart 1 starting
hart 2 starting
init: starting sh
$ uthread
t1:0x00000000000000D90,t20x00000000000002E08,t3:0x00000000000004E80,t4:0x0000000000006EF8
thread_a started
thread_b started
thread_c started
thread_c 0
thread_a 0
thread_b 0
thread_c 1
thread_a 1
thread_b 1
thread_c 2
thread_a 2
thread_b 2
thread_c 3
thread_a 3
thread_b 3
thread_c 4
thread_a 4
thread_b 4
thread_c 5
thread_a 5
thread_b 5
```

```
chenxinrui@chenxinrui-VirtualBox: ~/xv6-labs
thread_c 92
thread_a 92
thread_b 92
thread_c 93
thread_a 93
thread_b 93
thread_c 94
thread_a 94
thread_b 94
thread_c 95
thread_a 95
thread_b 95
thread_c 96
thread_a 96
thread_b 96
thread_c 97
thread_a 97
thread_b 97
thread_c 98
thread_a 98
thread_b 98
thread_c 99
thread_a 99
thread_b 99
thread_c: exit after 100
thread_a: exit after 100
thread_b: exit after 100
thread_schedule: no runnable threads
$
```

2. Xv6 lab: Lock/Memory allocator

第一步：在 `kalloc.c` 中首先将 `kmem` 修改为数组，这样每个 `cpu` 对应一份 `freelist` 和 `lock`

```
struct {
    struct spinlock lock;
    struct run *freelist;
} kmem[NCPU];
```

第二步：初始化 `kmem` 时将每个 `cpu` 对应 `kmem[i]` 都初始化

```
void
kinit()
{
    for(int i=0; i<NCPU; i++)
    {
        initlock(&kmem[i].lock, "kmem");
    }
    freerange(end, (void*)PHYSTOP);
}
```

第三步：修改 kfree 代码

```
void
kfree(void *pa)
{
    struct run *r;

    if(((uint64)pa % PGSIZE) != 0 || (char*)pa < end || (uint64)pa >=
PHYSTOP)
        panic("kfree");

    // Fill with junk to catch dangling refs.
    memset(pa, 1, PGSIZE);

    r = (struct run*)pa;

    push_off();// turn interrupts off
    int i=cuid();// core number
    acquire(&kmem[i].lock);
    r->next = kmem[i].freelist;
    kmem[i].freelist = r;
    release(&kmem[i].lock);

    pop_off();//turn interrupts on
}
```

第四步：修改 kalloc 代码

```
void *kalloc(void){
    struct run *r;
    push_off();// turn interrupts off
    int i=cuid();// core number
    acquire(&kmem[i].lock);
    r = kmem[i].freelist;
    if(r)
        kmem[i].freelist = r->next;
    release(&kmem[i].lock);
    if(!r){//current cpu->freelist is empty
        for(int j=0;j<NCPU;j++){//borrow from other cpu->freelist
            if(j!=i){
                acquire(&kmem[j].lock);
                if(kmem[j].freelist){
                    r=kmem[j].freelist;
                    kmem[j].freelist=r->next;
                    release(&kmem[j].lock);
                    break;
                }
                release(&kmem[j].lock);
            }
        }
    }
}
```

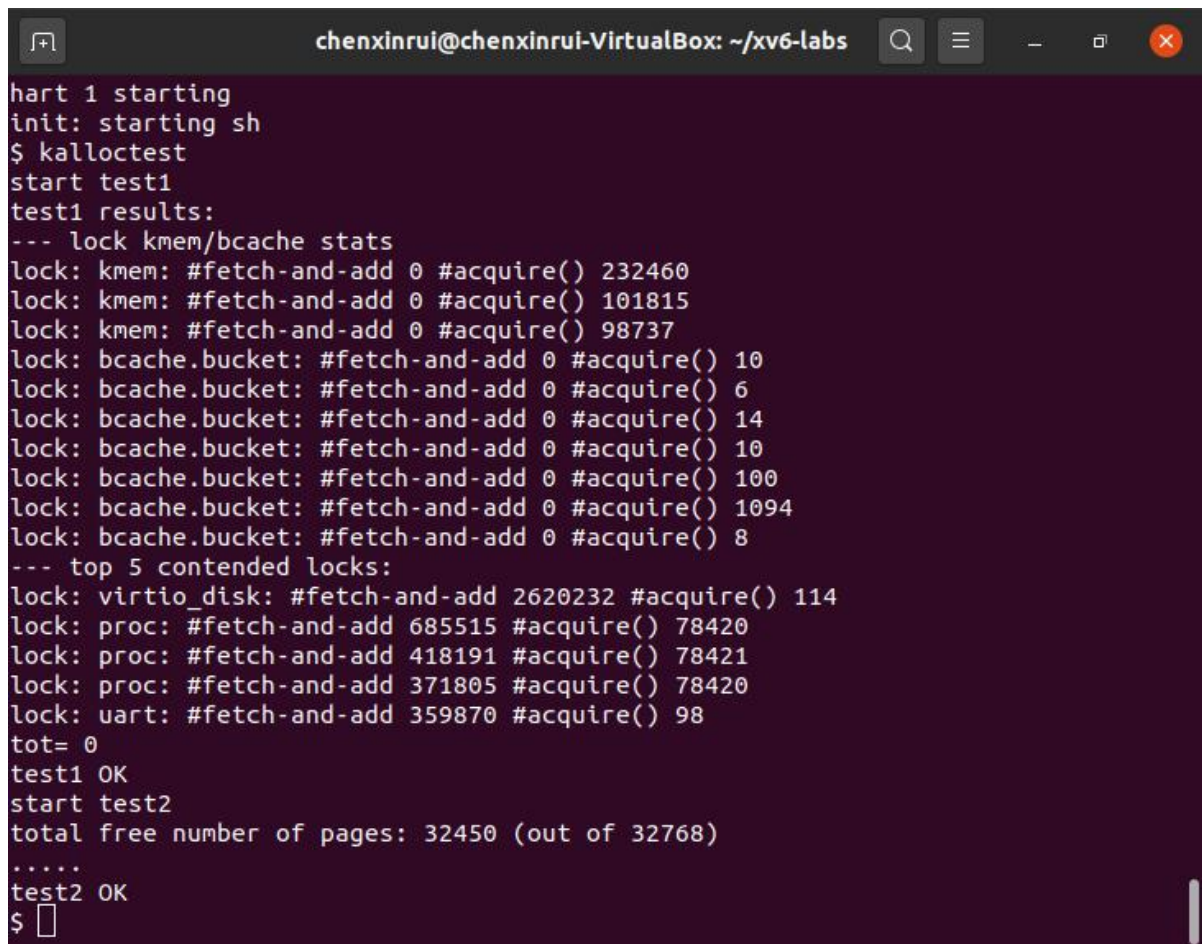
```

    }
}
}
pop_off();//turn on interrupt
if(r)
    memset((char*)r, 5, PGSIZE); // fill with junk
return (void*)r;
}

```

运行结果:

运行 `kalloctest` 以查看您的实现是否减少了锁争用, 并运行 `usertests sbrkmuch` 来检查它是否仍可以分配所有内存。



```

chenxinrui@chenxinrui-VirtualBox: ~/xv6-labs
hart 1 starting
init: starting sh
$ kalloctest
start test1
test1 results:
--- lock kmem/bcache stats
lock: kmem: #fetch-and-add 0 #acquire() 232460
lock: kmem: #fetch-and-add 0 #acquire() 101815
lock: kmem: #fetch-and-add 0 #acquire() 98737
lock: bcache.bucket: #fetch-and-add 0 #acquire() 10
lock: bcache.bucket: #fetch-and-add 0 #acquire() 6
lock: bcache.bucket: #fetch-and-add 0 #acquire() 14
lock: bcache.bucket: #fetch-and-add 0 #acquire() 10
lock: bcache.bucket: #fetch-and-add 0 #acquire() 100
lock: bcache.bucket: #fetch-and-add 0 #acquire() 1094
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8
--- top 5 contended locks:
lock: virtio_disk: #fetch-and-add 2620232 #acquire() 114
lock: proc: #fetch-and-add 685515 #acquire() 78420
lock: proc: #fetch-and-add 418191 #acquire() 78421
lock: proc: #fetch-and-add 371805 #acquire() 78420
lock: uart: #fetch-and-add 359870 #acquire() 98
tot= 0
test1 OK
start test2
total free number of pages: 32450 (out of 32768)
.....
test2 OK
$ 

```



```
chenxinrui@chenxinrui-VirtualBox: ~/xv6-labs
test1 results:
--- lock kmem/bcache stats
lock: kmem: #fetch-and-add 0 #acquire() 232460
lock: kmem: #fetch-and-add 0 #acquire() 101815
lock: kmem: #fetch-and-add 0 #acquire() 98737
lock: bcache.bucket: #fetch-and-add 0 #acquire() 10
lock: bcache.bucket: #fetch-and-add 0 #acquire() 6
lock: bcache.bucket: #fetch-and-add 0 #acquire() 14
lock: bcache.bucket: #fetch-and-add 0 #acquire() 10
lock: bcache.bucket: #fetch-and-add 0 #acquire() 100
lock: bcache.bucket: #fetch-and-add 0 #acquire() 1094
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8
--- top 5 contended locks:
lock: virtio_disk: #fetch-and-add 2620232 #acquire() 114
lock: proc: #fetch-and-add 685515 #acquire() 78420
lock: proc: #fetch-and-add 418191 #acquire() 78421
lock: proc: #fetch-and-add 371805 #acquire() 78420
lock: uart: #fetch-and-add 359870 #acquire() 98
tot= 0
test1 OK
start test2
total free number of pages: 32450 (out of 32768)
.....
test2 OK
$ usertests sbrkmuch
usertests starting
test sbrkmuch: OK
ALL TESTS PASSED
$
```

3. Xv6 lab: Lock/Buffer cache

第一步：在 kernel/buf.h 的 struct 中添加字段 time_stamp，用以标记 buf 的时间戳

```
struct buf {
    int valid;    // has data been read from disk?
    int disk;    // does disk "own" buf?
    uint dev;
    uint blockno;
    struct sleeplock lock;
    uint refcnt;
    struct buf *prev; // LRU cache list
    struct buf *next;
    uchar data[BSIZE];
    uint time_stamp;
};
```

第二步: bget()中可能出现两种情况: hit(命中), eviction (驱逐之前的替换新的)。

在这里将 bcache 设计成如下形式，其中 NBUCKET 需要小一点，因为锁的数目有限，而 binit 至少需要 $NBUF * NBUCKET + NBUCKET$ 个锁

```
#define NBUCKET 7

extern uint ticks;

struct {
    struct spinlock lock;
    struct buf buf[NBUF]; // NBUF=30

    // Linked list of all buffers, through prev/next.
    // Sorted by how recently the buffer was used.
    // head.next is most recent, head.prev is least.
    // struct buf head;

} bcache[NBUCKET];
```

第三步：修改后的 kernel/bio.c 如下

```
uint idx(uint blockno){
    return blockno % NBUCKET;
}

void binit(void){
    struct buf *b;
    for(int i=0;i<NBUCKET;i++){
        for(b = bcache[i].buf; b < bcache[i].buf+NBUF; b++){
            initsleeplock(&b->lock, "buffer");
        }
        initlock(&bcache[i].lock, "bcache.bucket");
    }
}

static struct buf* bget(uint dev, uint blockno){
    struct buf *b=0;
    int i= idx(blockno);
    uint min_time_stamp=-1;
    struct buf *min_b=0;
    acquire(&bcache[i].lock);//only need to hold one lock to avoid competition
    // Is the block already cached?
    for(b = bcache[i].buf; b < bcache[i].buf+NBUF; b++){
        if(b->dev==dev && b->blockno == blockno){//hit
            b->refcnt++;
```



```

        release(&bcache[i].lock);
        acquiresleep(&b->lock);
        return b;
    }//find the buf according to the smallest time_stamp
    if(b->refcnt==0 && b->time_stamp<min_time_stamp){
        min_time_stamp=b->time_stamp;
        min_b=b;
    }
}
b=min_b;
if(b!=0){
    b->dev = dev;
    b->blockno = blockno;
    b->valid = 0;
    b->refcnt = 1;
    release(&bcache[i].lock);
    acquiresleep(&b->lock);
    return b;
}
panic("bget: no buffers");}
void brelse(struct buf *b){
    if(!holdingsleep(&b->lock))
        panic("brelse");
    releasesleep(&b->lock);
    int i=idx(b->blockno);
    acquire(&bcache[i].lock);
    b->refcnt--;
    if(b->refcnt == 0){// no one is waiting for it.
        b->time_stamp=ticks;
    }
    release(&bcache[i].lock);}
void bpin(struct buf *b){
    int i=idx(b->blockno);
    acquire(&bcache[i].lock);
    b->refcnt++;
    release(&bcache[i].lock);
}
void bunpin(struct buf *b) {
    int i=idx(b->blockno);
    acquire(&bcache[i].lock);
    b->refcnt--;
    release(&bcache[i].lock);
}

```

运行结果：

```
chenxinrui@chenxinrui-VirtualBox: ~/xv6-labs
hart 1 starting
hart 2 starting
init: starting sh
$ bcachetest
start test0
test0 results:
--- lock kmem/bcache stats
lock: kmem: #fetch-and-add 0 #acquire() 32927
lock: kmem: #fetch-and-add 0 #acquire() 83
lock: kmem: #fetch-and-add 0 #acquire() 52
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8296
lock: bcache.bucket: #fetch-and-add 0 #acquire() 10298
lock: bcache.bucket: #fetch-and-add 0 #acquire() 10442
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8748
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8476
lock: bcache.bucket: #fetch-and-add 0 #acquire() 9538
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8288
--- top 5 contended locks:
lock: virtio_disk: #fetch-and-add 11429147 #acquire() 1014
lock: proc: #fetch-and-add 2847707 #acquire() 82846
lock: proc: #fetch-and-add 2010248 #acquire() 83177
lock: proc: #fetch-and-add 1948997 #acquire() 82825
lock: proc: #fetch-and-add 1822270 #acquire() 82824
tot= 0
test0: OK
start test1
test1 OK
$
```

4. Xv6 lab: File System/Large files

第一步：首先修改 fs.h 中的宏定义，NINDIRECT_1 是第一级间接访问的 block 数目，NINDIRECT_2 是第二级间接访问的 block 数目

```
#define FSMAGIC 0x10203040

#define NDIRECT 11
#define NINDIRECT_1 (BSIZE / sizeof(uint))
#define NINDIRECT_2 ( NINDIRECT_1 * NINDIRECT_1 )
#define NINDIRECT (NDIRECT + NINDIRECT_2)
#define MAXFILE (NDIRECT + NINDIRECT)
```

第二步: 修改 bmap () , 使其除了直接块和单间接块之外还实现双间接块。ip->addrs [] 的前 11 个元素应该是直接块; 第十二个应该是一个间接块 (就像当前的块一样); 第 13 个应该是一个新的双间接块。

```
static uint bmap(struct inode *ip, uint bn){
    uint addr, *a;
    struct buf *bp;
    if(bn < NDIRECT){
        if((addr = ip->addrs[bn]) == 0)
            ip->addrs[bn] = addr = balloc(ip->dev);
        return addr;
    }
    bn -= NDIRECT;
    if(bn < NINDIRECT_1){
        // Load first-level indirect block, allocating if necessary.
        if((addr = ip->addrs[NDIRECT]) == 0)
            ip->addrs[NDIRECT] = addr = balloc(ip->dev);
        bp = bread(ip->dev, addr);
        a = (uint*)bp->data;
        if((addr = a[bn]) == 0){
            a[bn] = addr = balloc(ip->dev);
            log_write(bp);
        }
        brelse(bp);
        return addr;
    }
    bn -= NINDIRECT_1;
    if(bn < NINDIRECT_2){
        // Load second-level indirect block, allocating if necessary.
        if((addr = ip->addrs[NDIRECT+1]) == 0)
            ip->addrs[NDIRECT+1] = addr = balloc(ip->dev);
        bp = bread(ip->dev, addr);
        a = (uint*)bp->data;
        uint bn_1 = (bn & 0xff00) >> 8; //Level1
        uint bn_2 = bn & 0xff; //Level2
        if((addr = a[bn_1]) == 0){
            a[bn_1] = addr = balloc(ip->dev);
            log_write(bp);
        }
        brelse(bp);
        bp = bread(ip->dev, addr);
        a = (uint*)bp->data;
```

```

    if((addr = a[bn_2]) == 0){
        a[bn_2] = addr = balloc(ip->dev);
        log_write(bp);
    }
    brelse(bp);
    return addr;
}
panic("bmap: out of range");
}

```

第三步：确保 itrunc 释放文件的所有块，包括双间接块。

```

void itrunc(struct inode *ip){
    int i, j, k;
    struct buf *bp;
    struct buf *bp2;
    uint *a;
    uint *a2;
    for(i = 0; i < NDIRECT; i++){
        if(ip->addrs[i]){
            bfree(ip->dev, ip->addrs[i]);
            ip->addrs[i] = 0;
        }
    }
    if(ip->addrs[NDIRECT]){
        bp = bread(ip->dev, ip->addrs[NDIRECT]);
        a = (uint*)bp->data;
        for(j = 0; j < NINDIRECT_1; j++){
            if(a[j])
                bfree(ip->dev, a[j]);
        }
        brelse(bp);
        bfree(ip->dev, ip->addrs[NDIRECT]);
        ip->addrs[NDIRECT] = 0;
    }
    if(ip->addrs[NDIRECT+1]){
        bp = bread(ip->dev, ip->addrs[NDIRECT+1]);
        a = (uint*)bp->data;
        for(j = 0; j < NINDIRECT_1; j++){
            if(a[j]){
                bp2 = bread(ip->dev, a[j]);
                a2 = (uint*)bp2->data;
                for(k=0 ;k < NINDIRECT_1 ;k++){
                    if(a2[k]) bfree(ip->dev,a2[k]);
                }
            }
        }
    }
}

```

